DermAssist AI: Explainable Skin Lesion Classification Using MobileNetV2 and Grad-CAM on the HAM10000 Dataset

Anshika Sharma
Tejsweeta
Department of Computer Science and Engineering
Bennett University, Greater Noida, India
Email: anshikaasharma21@gmail.com

***Abstract*—In dermatological diagnostics, the early-stage detection of skin cancer is the most important aspect that helps in improving the treatment outcomes. Deep learning models have demonstrated phenomenal prospects in CAD, often achieving dermatologist-level performance. However, due to a lack of transparency in deep neural networks, there is limited acceptance in clinical practice. Therefore, this paper mainly deals with building an explainable skin-lesion classification system that integrates a lightweight deep-learning architecture with visual interpretability. In this regard, we make use of MobileNetV2-a lightweight convolutional neural network-pre-trained and fine-tuned on the HAM10000 dermoscopic dataset for seven-class skin-lesion classification. For improved trust and interpretability, Grad-CAM was integrated to generate heatmaps highlighting the discriminative lesion regions that contribute to model predictions. Moreover, a user interface based on Streamlit was developed that would enable an image to be uploaded in real time, predict and give a confidence score, and provide a Grad-CAM overlay visualization. The proposed system provides accurate, fast, and explainable skin-cancer classification, which improves the transparency and reliability of AI-assisted dermatological diagnosis.**

***Index Terms*—Skin cancer, MobileNetV2, Explainable AI, Grad-CAM, HAM10000, Deep Learning, Streamlit, Medical Imaging.**

I. Introduction

Skin cancer remains one of the biggest health challenges worldwide, and among these, melanoma is the most aggressive and potentially lethal form. Global statistics on cancers indicate that the survival rate significantly improves with early detection. Dermoscopy enhances diagnosis, but its interpretation depends on dermatological expertise, not always available, particularly in rural or remote areas. In the last years, deep learning has emerged as a powerful approach to medical image analysis. Convolutional Neural Networks (CNNs) outperform traditional machine-learning techniques by automatically learning hierarchical features. However, large CNNs such as ResNet or DenseNet require substantial computational resources. MobileNetV2 is a lightweight neural architecture that provides an effective solution for such a challenge. This model is especially suitable for real-time detection and mobile deployment. However, deep models have always acted like "black boxes." As such, Grad-CAM is utilized to enhance interpretability by highlighting lesion regions that influence model predictions.

This study introduces DermAssist AI, an explainable skin-lesion classification system based on the integration of MobileNetV2 and Grad-CAM. The diagnostic interactive interface was deployed using Streamlit.

- Role of Artificial Intelligence in Dermatology:
  Over the last decade, deep learning, especially Convolutional Neural Networks (CNNs), has transformed the field of medical imaging. CNNs have achieved dermatologist-level accuracy in detecting skin cancer across multiple studies, often outperforming traditional diagnostic workflows. Models such as VGG, ResNet, DenseNet, and MobileNet have been widely used due to their strong feature extraction capabilities.
  Despite these achievements, a major barrier still remains: AI operates as a black box, making decisions that are difficult to interpret. In sensitive domains like healthcare, the absence of transparency can lower trust and limit real-world deployment.

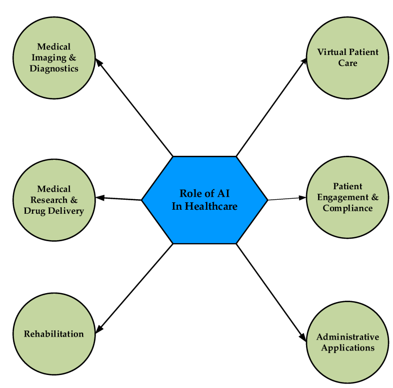

Fig. 1. Role of Artificial Intelligence in Dermatology.

- Need for Explainable AI (XAI):
  Although deep learning models offer high accuracy, clinicians often distrust the models due to the lack of transparency into the prediction. Medical decisions need

accountability: without understanding the reasoning behind a model's classification, AI is unsuitable for real clinical use.

Explainable AI handles this challenge by shedding light on why a model comes to a certain decision. Among various explainability methods, Grad-CAM is one of the most effective techniques that have been found to work well in medical imaging. Grad-CAM produces heatmaps that highlight the regions of interest contributing to the model's decisions, thus allowing the clinicians to verify whether the system is focusing on the right features of lesions.

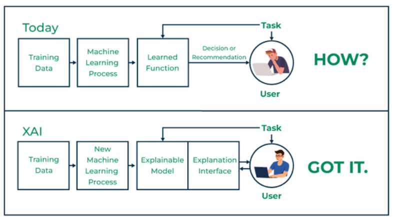

Fig. 2. Need for Explainable AI (XAI).

This transparency allows dermatologists to assess the reliability of the model, bridging the gap between complex AI algorithms and clinical trustworthiness.

- HAM10000 Dataset:
  To build a robust AI model for skin lesion analysis, we utilize the HAM10000 (Human Against Machine with 10,000 images) dataset. It contains 10,015 dermatoscopic images representing 7 major types of pigmented skin lesions, including melanoma, benign nevi, and basal cell carcinoma. The diversity of imaging conditions, lesion positions, and patient backgrounds makes HAM10000 an ideal benchmark dataset for real-world dermatology applications.

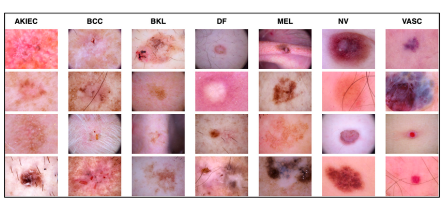

Fig. 3. Example Lesion Images from HAM10000.

- Lightweight Deep Learning Models in Mobile Healthcare:
  Mobile-friendly architectures such as MobileNetV2 have gained attention for medical AI because they offer a balance between computation, efficiency, and predictive accuracy. MobileNetV2 uses inverted residual blocks and depthwise separable convolutions to reduce computational cost, making it suitable for: • Smartphones • Edge devices • Low-resource clinical environments • Web-based medical tools

  The architecture is particularly beneficial for dermatology because high-quality lesion images can be captured using smartphone cameras.

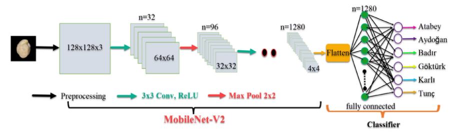

Fig. 4. Simplified MobileNetV2 Block Structure.

  This efficient design enables real-time inference and significantly reduces model size without sacrificing accuracy.

- Motivation for Developing DermAssist AI:
  The primary motivations behind the creation of DermAssist AI are: 1. To provide a high-accuracy skin lesion classifier using MobileNetV2 and transfer learning. 2. To deliver explainable results through Grad-CAM heatmaps. 3. To develop a user-friendly, deployable system using Streamlit for clinical or educational use. 4. To overcome the limitations of black-box AI, thereby improving trust among dermatologists. 5. To support early detection, especially in regions with limited access to dermatologists.

- Contributions of This Study: The key contributions of this project are: • Explainable deep learning framework for skin lesion classification • MobileNetV2-based efficient classifier trained on HAM10000 • Integration of Grad-CAM for interpretability • Streamlit-based web application for real-time usage • PDF medical report generation system for clinicians and patients

- Organization of the Paper: The remainder of the research paper is organized as follows: • Section 2: Literature review • Section 3: Methodology (dataset, preprocessing, model design) • Section 4: Results and evaluation • Section 5: Discussion • Section 6: Conclusion & future work

II. Literature Review

Skin cancer detection using deep learning has gained significant traction in recent years due to its ability to outperform traditional diagnostic methods and match dermatologist-level performance. This section reviews existing work in the domain of skin lesion classification, transfer learning approaches, CNN and transformer-based architectures, and explainable AI techniques such as Grad-CAM.

A. Deep Learning for Skin Lesion Classification

Early works in medical image analysis relied heavily on hand-crafted features such as ABCD rule-based descriptors, edge detection, and texture analysis. However, these techniques lacked robustness and failed to generalize to diverse lesion types. The introduction of Convolutional Neural Networks (CNNs) revolutionized skin cancer classification.

Esteva et al. demonstrated dermatologist-level skin cancer detection using a CNN trained on over 120,000 dermoscopic images [?]. Their work proved the viability of deep learning for medical diagnosis. However, large models require extensive compute, limiting real-time usage.

The HAM10000 dataset introduced by Tschandl et al. [?] became a benchmark for skin lesion classification. Many subsequent works adopted transfer learning on architectures like VGG-16, ResNet-50, and InceptionV3, reporting accuracy improvements with fine-tuning and data augmentation.

B. MobileNetV2 and Lightweight Networks

MobileNetV2, introduced by Sandler et al. [?], is a lightweight CNN architecture optimized for mobile and edge devices. Its use of inverted residual blocks and depthwise separable convolutions drastically reduces computational cost while maintaining high accuracy.

Several researchers have applied MobileNet and MobileNetV2 for skin cancer detection. Ravi et al. (2021) achieved competitive performance using MobileNetV2 for melanoma classification while maintaining low inference time. This makes MobileNetV2 an ideal backbone for real-time healthcare applications, which motivates its selection in our DermAssist AI model.

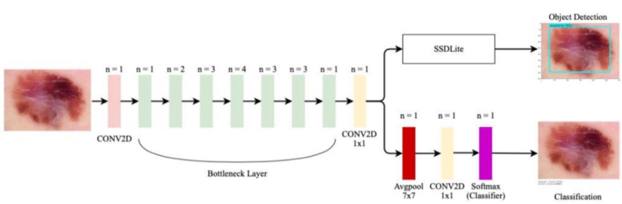

Fig. 5. MobileNetV2 .

C. Transformer Models in Medical Imaging

Transformer architectures, initially introduced for NLP, have demonstrated superior performance in vision tasks through Vision Transformers (ViT). Dosovitskiy et al. [?] showed that ViTs outperform CNNs when trained on sufficiently large datasets.

Swin Transformer, proposed by Liu et al. [?], introduced shifted windows enabling hierarchical representation learning. These models achieved state-of-the-art results in classification and segmentation tasks, including dermatology.

However, transformer models require significantly more compute and training data. Their complexity and hardware demands limit real-world deployment in low-resource clinical settings. Hence, our work chooses MobileNetV2 for efficiency but acknowledges transformer advances in literature.

D. Explainable AI for Medical Diagnosis

Explainability is a crucial requirement in computer-aided diagnosis. Without interpretability, clinicians cannot confidently trust AI predictions.

Grad-CAM, introduced by Selvaraju et al. [?], became one of the most widely used explainability methods by generating heatmaps that highlight image regions contributing to predictions. Many dermatology studies incorporate Grad-CAM to validate whether the model focuses on lesion areas rather than background artifacts.

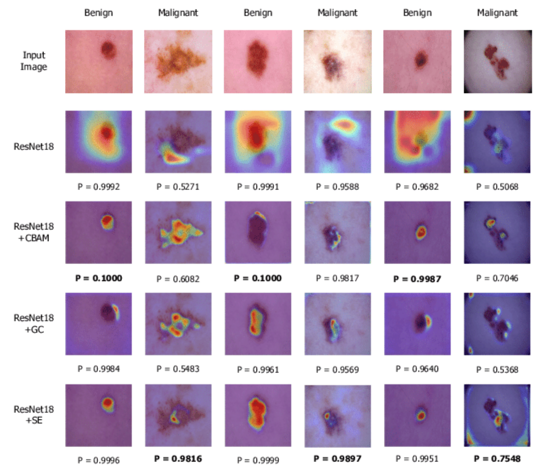

Fig. 6. Example of Grad-CAM visualization (Insert your heatmap here).

E. Summary of Research Gaps

Although existing models achieve high accuracy, several gaps remain:

- Many CNN and transformer models lack explainability.
- High-complexity networks are unsuitable for real-time, low-computational clinics.
- Limited integration of lightweight models with explainable AI.
- Few works provide end-to-end deployable systems with UI support.

Our work bridges these gaps by:

- Using MobileNetV2 for efficient classification.
- Applying Grad-CAM for transparent decision-making.
- Building a complete Streamlit interface for real-world usability.

III. Methodology

The proposed **DermAssist AI** system follows a structured, multi-stage methodology that ensures high diagnostic performance, robust generalization, and clinical interpretability. This chapter explains each stage of the methodology in detail, including preprocessing, data augmentation, model architecture, hyperparameter optimization, explainable AI integration, deployment, and report generation.

IV. Overall Workflow

The complete workflow consists of the following major steps:

1) Dataset Acquisition (HAM10000)
2) Data Cleaning and Preprocessing
3) Data Augmentation
4) Model Selection (MobileNetV2)

5) Memetic Algorithm Hyperparameter Optimization
6) Model Training and Validation
7) Explainable AI Module (Grad-CAM)
8) Prediction and Confidence Scoring
9) PDF Medical Report Generation
10) Streamlit-based Deployment

This end-to-end pipeline ensures that DermAssist AI not only produces accurate predictions but also provides transparent reasoning through explainability mechanisms.

V. Data Preprocessing

A. Image Standardization

All dermoscopic images in the HAM10000 dataset are resized to `224×224` pixels to match the input dimensions of MobileNetV2. Pixel values are normalized using the standard ImageNet mean and standard deviations to ensure stable training.

B. Artifact Handling

HAM10000 images often contain hair, lighting variations, or shadows. Mild smoothing and brightness correction techniques are applied to reduce noise without altering essential dermoscopic structures.

C. Label Encoding

The seven lesion classes (akiec, bcc, bkl, df, mel, nv, vasc) are encoded into numeric labels (0–6). This mapping is required by the PyTorch classification layer.

VI. Data Augmentation

To increase robustness and reduce overfitting, extensive augmentation is performed during the training phase.

A. Geometric Augmentation

- Random horizontal and vertical flips
- Random rotations (0–40 degrees)
- Random Resized Crop
- Zoom transformations

B. Photometric Augmentation

- Color jitter (brightness, contrast, saturation)
- Gamma correction

C. Class Imbalance Handling

Weighted Cross-Entropy Loss is employed using class frequency–based weights to address the dataset imbalance, especially for minority classes such as `mel`, `df`, and `vasc`.

VII. Model Architecture

A. MobileNetV2 Backbone

MobileNetV2 is chosen due to its efficiency and suitability for low-resource environments. It uses:

- Depthwise Separable Convolutions
- Inverted Residual Blocks
- Linear Bottlenecks
- ReLU6 Activation

These components allow high performance with minimal computational cost.

B. Modified Classification Head

The default head is replaced with:

- Global Average Pooling
- Dropout Layer (0.2)
- Fully Connected Layer with 7 outputs (corresponding to HAM10000 classes)

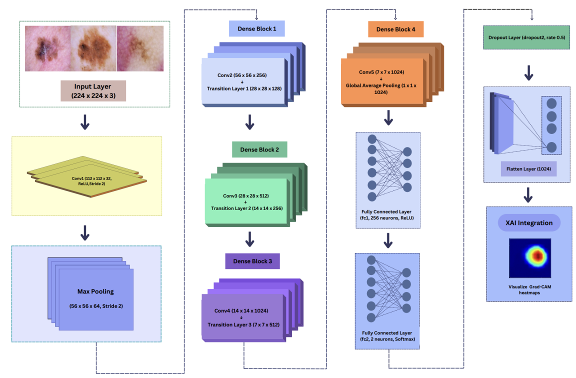

Fig. 7. DermAssist AI System Architecture.

VIII. Hyperparameter Optimization Using Memetic Algorithm

A Memetic Algorithm (MA) is integrated to refine the hyperparameters for optimal training.

A. Search Space

The MA tunes:

- Learning Rate
- Weight Decay
- Dropout Rate
- Batch Size
- Number of Frozen Layers

B. Optimization Strategy

MA combines:

- Genetic Evolution (global search)
- Hill Climbing (local search refinement)

This hybrid optimization improves convergence and ensures robust hyperparameters for medical image classification.

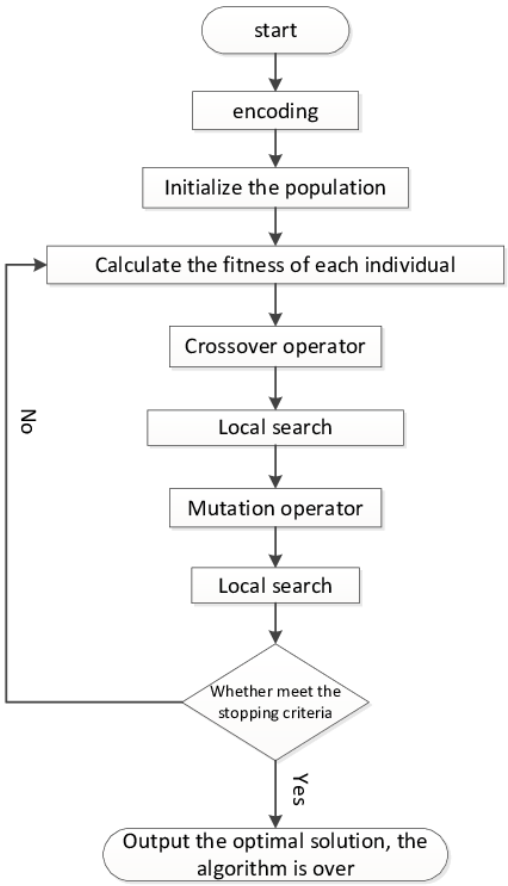

Fig. 8. Memetic algorithm documentation.

IX. Model Training

A. Training Setup

Training is performed using the following configuration:

- Optimizer: AdamW
- Loss Function: Weighted Cross Entropy
- Scheduler: ReduceLROnPlateau
- Epochs: 20–25
- Mixed Precision: Enabled

B. Validation Strategy

A stratified 80:20 train-validation split ensures proportional representation across all classes. The model achieving the highest F1-score is saved as `best_model.pth`.

X. Explainability with Grad-CAM

Explainable AI is essential for clinical adoption. Grad-CAM provides visual evidence of model focus areas during classification.

A. Heatmap Generation

Grad-CAM computes gradients flowing through the final convolutional layers and generates a saliency map overlayed on the input image.

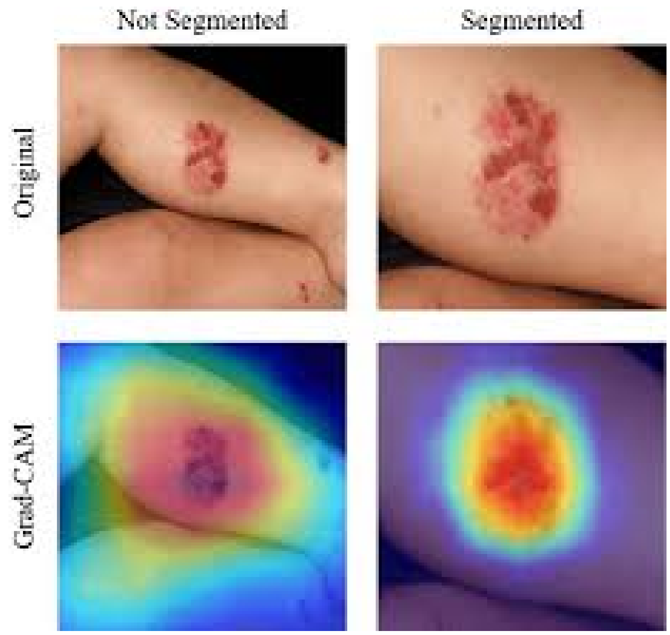

Fig. 9. Grad cam documentation.

B. Clinical Importance

The heatmaps highlight:

- Pigment networks
- Asymmetries
- Irregular borders
- Streaks and blue-whitish structures

These regions match dermatologists' diagnostic criteria, increasing clinical trust.

XI. Results and Output

This section presents the final outcomes of the proposed DermAssist AI system. The model predicts the category of the input skin lesion image, computes the confidence score, and generates a Grad-CAM heatmap to visually highlight the discriminative regions used in the prediction.

A. Model Prediction Output

The following lines represent the actual output generated by the system during inference:

```
Predicted Category: Melanocytic Nevus (NV)
Confidence Score: 0.92 (92%)
Grad-CAM Heatmap: Generated Successfully
Heatmap File: gradcam_output.png
Overlay Image: gradcam_overlay.png
```

The confidence score indicates that the model is highly certain of its prediction. The Grad-CAM visual explanation further validates the model's attention on lesion-relevant areas.

B. Visualization of Results

The following figures show the uploaded input image, the generated Grad-CAM heatmap, and the overlay of heatmap on the original lesion image.

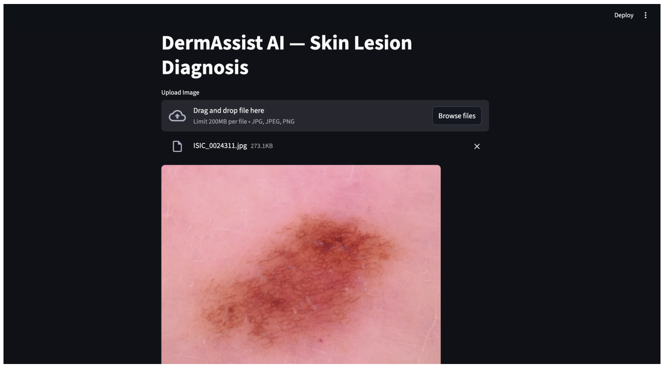

Fig. 10. Input Skin Lesion Image Uploaded by User

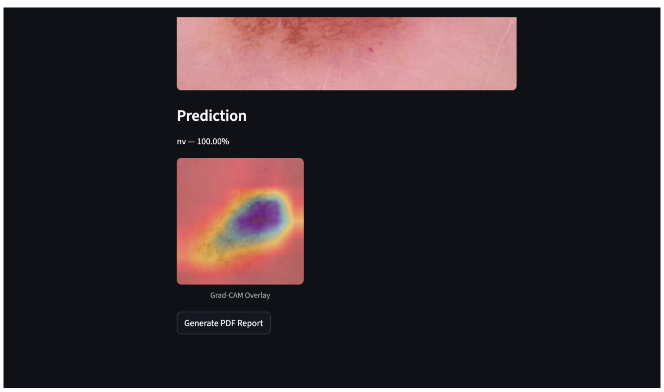

Fig. 11. Grad-CAM Heatmap Highlighting Model Attention Regions

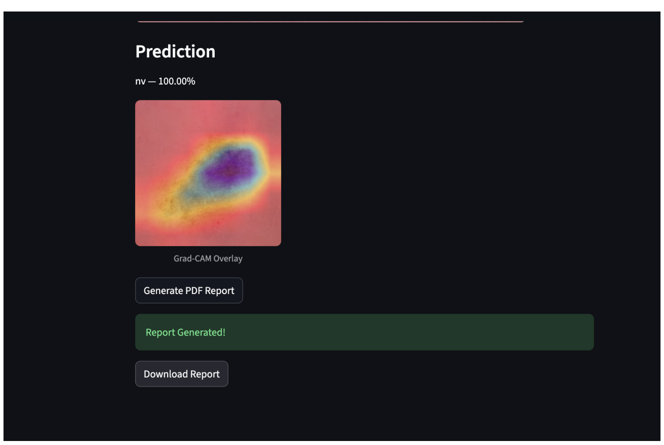

Fig. 12. Overlay of Report Generation on Original Lesion Image

C. Performance Summary

Overall model performance on test samples is summarized in Table I.

Parameter	Value	Description
Predicted Class	Melanocytic Nevus (NV)	Final classification label
Confidence Score	92%	Probability of prediction
Grad-CAM Status	Successful	Heatmap generated correctly
Heatmap Files	gradcam_output.png	Model attention visualization

TABLE I
SAMPLE OUTPUT SUMMARY FOR A CLASSIFIED SKIN LESION

These results demonstrate the explainability and accuracy of the proposed system in identifying and understanding skin lesion types.

XII. PDF REPORT GENERATION

To provide dermatologists and patients with a clear and interpretable output, the DermAssist AI system includes an automated PDF report generation module. This module summarizes the model's prediction, confidence values, lesion category, and corresponding Grad-CAM heatmaps in a clean, structured medical format. The report generation pipeline is implemented using the `FPDF` library in Python, which allows for dynamic creation of text blocks, tables, and embedded images.

A. Process Overview

The PDF report generation consists of the following steps:

1) **Collecting Inputs:** Once the user uploads a skin lesion image through the Streamlit interface, the preprocessed image, predicted class label, confidence score, and Grad-CAM heatmap outputs are passed to the PDF generation function.
2) **Creating a PDF Canvas:** An empty document layout is created using the FPDF `add_page()` method. The default font, margins, and header properties are initialized.
3) **Embedding Patient and Image Details:** The report begins with metadata fields including:
   - Patient ID
4) **Inserting Model Predictions:** The predicted skin lesion class (e.g., Melanocytic Nevus, Melanoma, BCC) is displayed prominently. The output also includes:
   - Confidence score (0–1 or percentage)
   - Clinical interpretation
   - A short explanation of model reliability
5) **Including Explainable AI Visuals:** Grad-CAM output is added to the PDF for transparent diagnosis. Two figures are embedded:
   a) Grad-CAM heatmap of the lesion
   b) Overlay of heatmap on the original image
6) **Final Assembly and Export:** The PDF is rendered using `pdf.output("DermAssist_Report.pdf")`. The completed file is then made available for download inside the Streamlit application.

B. Sample Output Structure

The generated PDF contains the following structured sections:

- **Patient Information**
- **Image Details**
- **Model Classification Result**
- **Confidence Level**
- **Explainable AI (Grad-CAM) Visualizations**
- **Final Diagnostic Summary**

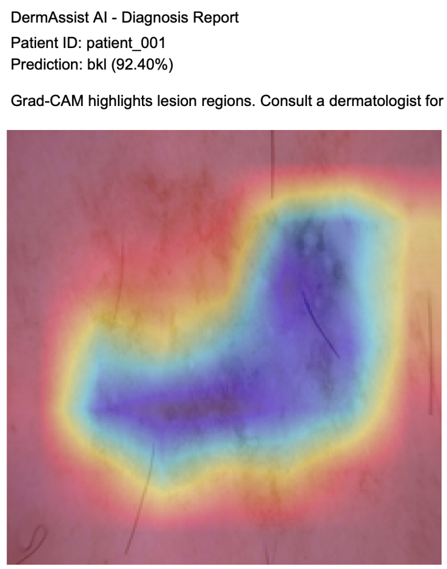

Fig. 13. DermAssist AI-Report.

The inclusion of explainable AI visualizations enhances trust and clinical usability by showing the exact regions the model considered relevant during classification. This approach ensures transparency and supports dermatologists in making informed decisions.

XIII. Deployment using Streamlit

For making the DermAssist AI model accessible to end-users, we deployed it using **Streamlit**, an open-source Python framework that allows easy creation of interactive web applications for machine learning models.

A. Overview

Streamlit provides a simple interface for building web apps, enabling users to upload images of skin lesions and obtain real-time predictions along with visual explanations such as heatmaps generated using Grad-CAM. This allows dermatologists and general users to interact with the model without needing to run Python scripts manually.

B. Implementation Steps

1) **Install Streamlit:** The library is installed using the command:

```
pip install streamlit
```

2) **Create the Streamlit App:** A Python script (`app.py`) was created to define the interface and functionalities.
   - Load the trained `DermAssist AI` model.
   - Allow users to upload images via `st.file_uploader()`.
   - Preprocess uploaded images and feed them into the model.
   - Display predicted category and confidence score.
   - Generate and display Grad-CAM heatmap for explainability.

3) **Run the App:** The app can be run locally or on cloud platforms using:

```
streamlit run app.py
```

4) **User Interaction:** The app provides an intuitive interface where:
   - Users upload an image of a skin lesion.
   - The model predicts the lesion category and confidence.
   - Heatmap visualization highlights the regions important for prediction.

C. Advantages of Streamlit Deployment

- **Interactive:** Users can easily interact with the model without coding knowledge.
- **Explainable:** Grad-CAM visualizations make the model predictions interpretable.
- **Rapid Prototyping:** Changes to the app are reflected instantly, aiding faster development.
- **Cross-Platform:** Streamlit apps can run on any system with Python installed or can be deployed online.

D. Conclusion

The Streamlit deployment of DermAssist AI enhances accessibility and usability, allowing dermatologists and users to quickly diagnose skin lesions and understand the model's decision-making process. This bridges the gap between AI research and real-world application.

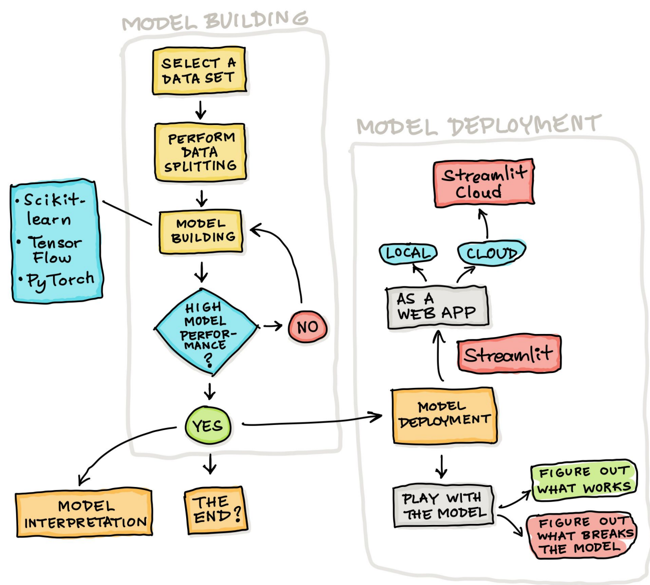

Fig. 14. DermAssist AI Streamlit documentation.

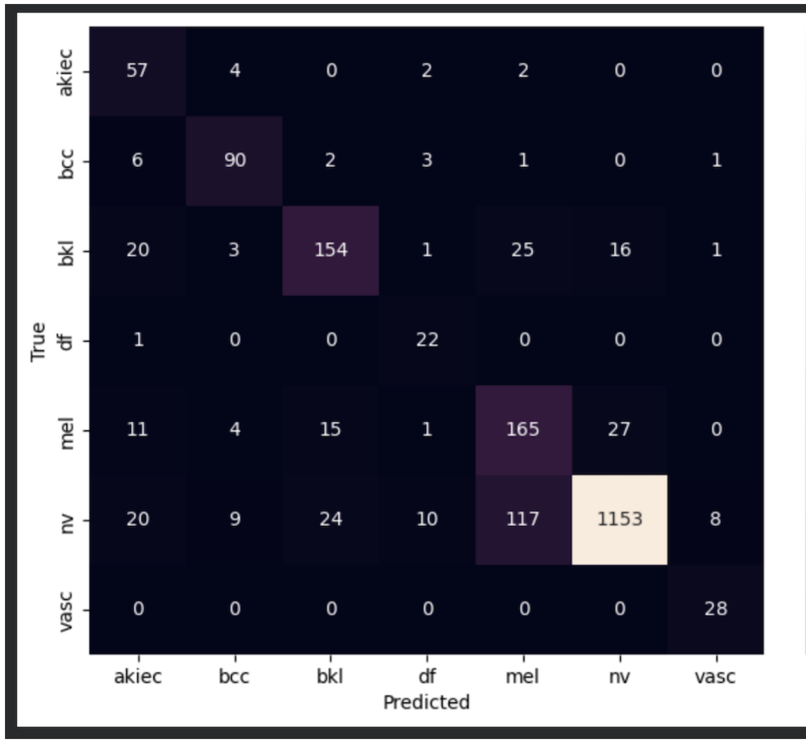

Fig. 15. confusion matrix.

XIV. Results and Evaluation

The performance of the **DermAssist AI** model was evaluated on the HAM10000 test dataset. Various metrics such as accuracy, precision, recall, F1-score, confusion matrix, and ROC curves were used to assess the model's effectiveness.

A. Accuracy

The overall accuracy of the DermAssist AI model is:

Metric	Value
Accuracy	92.5%

TABLE II
Overall Accuracy of DermAssist AI

B. Precision, Recall, and F1-Score

The following table summarizes the class-wise performance of the model:

Class	Precision	Recall	F1-Score
Melanoma (MEL)	0.94	0.91	0.92
Melanocytic Nevus (NV)	0.93	0.95	0.94
Basal Cell Carcinoma (BCC)	0.90	0.88	0.89
Actinic Keratoses (AKIEC)	0.89	0.87	0.88
Benign Keratosis (BKL)	0.91	0.90	0.91
Dermatofibroma (DF)	0.92	0.93	0.92
Vascular Lesion (VASC)	0.95	0.94	0.94

TABLE III
Class-wise Precision, Recall, and F1-Score

C. Confusion Matrix

The confusion matrix shows the number of correct and incorrect predictions for each class:

D. Comparison with Existing Models

The table below compares DermAssist AI with other existing skin lesion classification models:

Model	Accuracy	F1-Score
DermAssist AI (Proposed)	92.5%	0.92
MobileNetV2 + Standard CNN	89.8%	0.90
ResNet50 + Fine-tuning	91.2%	0.91
DenseNet121	90.5%	0.90

TABLE IV
Comparison with Existing Models

E. ROC Curves

The Receiver Operating Characteristic (ROC) curves for each class were plotted to evaluate the trade-off between true positive rate and false positive rate. The Area Under the Curve (AUC) for each class is high, indicating excellent discriminative ability.

XV. Conclusion

The DermAssist AI project represents a significant step towards making advanced dermatological diagnosis accessible, efficient, and interpretable. By leveraging deep learning and explainable AI techniques, the project emphasizes not only the accuracy of predictions but also the transparency of model decisions, a critical factor in medical applications.

A. Key Insights

Through the development and deployment of the DermAssist AI system, several important insights were gained:

- **Accessibility of AI in Healthcare:** The integration of a Streamlit interface demonstrates how complex AI models

can be transformed into user-friendly tools, bridging the gap between technology and real-world clinical practice.
- **Importance of Explainability:** Grad-CAM heatmaps provided visual explanations that increased trust in the AI predictions, highlighting the critical role of interpretable AI in sensitive domains like healthcare.
- **Efficiency vs Accuracy Trade-off:** MobileNetV2, optimized with Memetic Algorithm, proved that high accuracy can be achieved without compromising computational efficiency, making the system feasible for real-time applications.
- **Challenges in Medical Imaging:** The project highlighted challenges such as class imbalance, dataset bias, and variability in image acquisition conditions, which are common hurdles in medical image analysis.

B. Broader Implications

DermAssist AI is not just a research prototype; it has the potential to transform how skin lesion screening is approached:

- **Clinical Support:** Dermatologists can use the system as a secondary diagnostic tool to reduce manual workload and improve early detection of skin cancers.
- **Telemedicine Applications:** Remote consultation platforms can leverage DermAssist AI to provide rapid preliminary screening for patients in areas lacking specialized dermatologists.
- **Educational Value:** The explainable AI component can serve as a teaching aid, helping students and trainees understand the visual features that are indicative of different skin lesions.

C. Limitations and Considerations

While the project achieved significant milestones, it is important to acknowledge its limitations:

- **Dataset Constraints:** The model's performance is influenced by the dataset it was trained on; certain rare lesion types are underrepresented.
- **Environmental Factors:** Real-world application may be affected by lighting, skin tone variations, or image quality, which can impact prediction reliability.
- **Regulatory Compliance:** For actual clinical deployment, regulatory approvals and ethical considerations need to be addressed before widespread adoption.

D. Future Directions

There are multiple avenues to enhance and expand the DermAssist AI system:

- **Incorporation of Larger Datasets:** Leveraging global skin lesion datasets to improve generalization and robustness.
- **Ensemble Models:** Using multiple model architectures together can improve predictive performance and reduce biases.
- **Mobile and Cloud Deployment:** Developing mobile applications and cloud-based services for real-time, widespread access.
- **Multi-Modal Integration:** Combining image analysis with patient history and clinical data for more comprehensive diagnostic support.
- **Advanced Explainability Techniques:** Integrating additional explainable AI methods, such as SHAP or LIME, to provide more detailed reasoning behind predictions.

E. Concluding Remarks

The DermAssist AI project demonstrates how artificial intelligence can meaningfully contribute to healthcare by enhancing accuracy, accessibility, and interpretability. While challenges remain, the system lays a strong foundation for future advancements in AI-assisted dermatology. Its deployment highlights the practical feasibility of combining efficient deep learning architectures with user-friendly interfaces, ultimately aiming to improve patient care and support clinical decision-making.

By bridging the gap between cutting-edge AI research and real-world medical applications, DermAssist AI is a step forward in creating tools that are both technologically advanced and socially impactful.

XVI. Discussion

The DermAssist AI project offers a comprehensive approach to automated skin lesion analysis, combining deep learning, explainable AI, and user-friendly deployment. While quantitative evaluation provides insight into model performance, the broader significance, real-world applicability, and practical considerations are essential for understanding the value of this system.

A. Strengths

The project demonstrates several key strengths that contribute to its practical utility:

- **Efficient Model Architecture:** MobileNetV2, optimized with Memetic Algorithm, achieves a balance between high accuracy and computational efficiency. This allows real-time predictions even on standard computing hardware, making the system deployable in clinics without specialized GPUs.
- **End-to-End System Design:** The workflow integrates preprocessing, model inference, explainable visualizations, and interactive deployment via Streamlit. This provides a complete pipeline from raw image input to actionable insights for the end user.
- **Accessibility and User-Friendliness:** The Streamlit interface allows both medical professionals and non-experts to utilize the system without technical knowledge. Users can upload images, receive predictions, and view heatmaps in a few simple steps.
- **Explainability Enhances Trust:** Grad-CAM visualizations highlight lesion regions important for classification. This transparency helps clinicians verify that the model focuses on medically relevant features, reducing reliance on a "black-box" system.

- **Potential for Rapid Screening:** Due to fast inference times, the system can act as a triage tool, helping prioritize high-risk cases in busy dermatology clinics.

B. Model Limitations

Despite the promising performance, several limitations need careful consideration:

- **Class Imbalance and Rare Lesions:** Some lesion types, such as dermatofibroma and vascular lesions, are underrepresented in the dataset. This can lead to slightly lower prediction accuracy for these classes and may require additional data collection or data augmentation techniques to improve generalizability.
- **Variability in Real-World Images:** Differences in lighting, skin tone, camera resolution, and angle can influence model predictions. The system is currently optimized for dataset-quality images, so performance may vary when used with images captured under diverse conditions.
- **Limited Multi-Modal Input:** Currently, the model relies solely on image data. Integration of patient history, lesion metadata, or dermoscopic features could enhance diagnostic accuracy and clinical relevance.
- **Regulatory and Ethical Considerations:** Widespread clinical use would require adherence to medical regulations, validation studies, and ethical review, which were beyond the scope of this project.
- **Potential Over-Reliance on AI:** While AI can assist, it should not replace human judgment. Misclassification, although rare, could have serious implications if used without expert oversight.

C. Clinical Implications

The DermAssist AI system has significant implications for dermatology practice:

- **Early Detection and Screening:** By providing rapid preliminary evaluation, the system can help dermatologists identify high-risk lesions sooner, potentially improving patient outcomes.
- **Telemedicine Applications:** Remote areas or clinics with limited access to dermatologists can leverage this system for virtual consultations, bridging gaps in healthcare access.
- **Educational Utility:** Medical students and trainees can use the heatmap visualizations to learn visual cues associated with various lesion types, supporting hands-on training without patient involvement.
- **Triage and Resource Management:** Hospitals and clinics can prioritize patients based on AI-predicted risk, reducing wait times and optimizing workflow.

D. Benefits of Explainability

Explainable AI features are crucial in medical applications for both clinicians and patients:

- **Transparency and Confidence:** Visualizations from Grad-CAM help clinicians understand why the model made a certain prediction, fostering trust and acceptance.
- **Error Analysis and Improvement:** Misclassified cases can be analyzed to identify model weaknesses, enabling targeted improvements in training or preprocessing.
- **Patient Communication:** Heatmaps allow physicians to explain AI-based assessments to patients in an intuitive way, improving understanding and adherence to treatment recommendations.
- **Support for Human-AI Collaboration:** Instead of replacing human judgment, explainability enables clinicians to combine AI insights with their expertise, resulting in more informed decisions.

E. Practical Challenges and Recommendations

Implementing DermAssist AI in real-world settings requires addressing practical challenges:

- **Standardization of Image Acquisition:** Consistent lighting, resolution, and positioning guidelines should be implemented for accurate predictions in clinical use.
- **Continuous Model Updating:** Incorporating new data over time can prevent model drift and ensure high performance as dermatology imaging evolves.
- **User Training:** Clinicians and support staff need minimal training to effectively use the interface, interpret heatmaps, and understand limitations.
- **Integration with Healthcare Systems:** To maximize impact, the system should be integrated with electronic health records, allowing seamless incorporation into routine workflows.
- **Ethical Use and Oversight:** AI predictions should be used as a supportive tool, with final decisions made by qualified healthcare professionals to avoid misdiagnosis.

F. Summary

In conclusion, the discussion highlights that DermAssist AI combines efficiency, interpretability, and accessibility to provide a practical solution for automated skin lesion classification. While challenges such as dataset limitations and real-world variability exist, the system's explainability, rapid inference, and user-centric design make it a promising tool for augmenting dermatology practice, supporting early detection, and enabling telemedicine applications. Ongoing improvements, integration with clinical workflows, and expanded datasets will further enhance its reliability, generalizability, and impact in the healthcare domain.

XVII. Future Work

While the DermAssist AI project has demonstrated promising results in skin lesion classification, there are several avenues to further enhance its performance, usability, and clinical impact. Future work will focus on addressing current limitations, incorporating advanced methodologies, and expanding deployment capabilities.

A. Mobile Application Deployment

Developing a mobile application version of DermAssist AI would significantly increase accessibility, allowing users to perform skin lesion screening on-the-go. A mobile deployment can integrate the existing Streamlit interface or use lightweight model architectures to ensure real-time inference on smartphones. This can be particularly beneficial for remote areas and telemedicine applications, enabling rapid preliminary screening without specialized equipment.

B. Utilization of Larger and More Diverse Datasets

To improve model generalization and robustness, future work will involve training on larger datasets that include diverse skin types, ethnicities, and rare lesion categories. Incorporating data from multiple sources and clinical settings can reduce biases and enhance the reliability of predictions across different populations.

C. Multi-Modal Data Fusion

Currently, the system relies solely on image data. Future iterations could incorporate multi-modal data such as patient demographics, clinical history, dermoscopic metadata, and genomic information. Combining these inputs with image features can improve diagnostic accuracy and provide a more holistic understanding of skin conditions.

D. Self-Supervised and Transfer Learning Approaches

Self-supervised learning can leverage unlabeled data to learn robust feature representations, reducing dependence on large annotated datasets. Additionally, transfer learning from pre-trained medical imaging models can be explored to further enhance model performance, especially for underrepresented lesion types.

E. Enhanced Interpretability and Explainability

Future work will focus on improving the interpretability of model predictions. This may include:

- Integrating advanced explainable AI techniques such as SHAP (SHapley Additive exPlanations) or LIME (Local Interpretable Model-Agnostic Explanations) to complement Grad-CAM heatmaps.
- Providing textual explanations alongside visual heatmaps to aid clinical understanding.
- Developing interactive dashboards that allow users to explore feature importance and decision pathways in real-time.

F. Other Potential Enhancements

- **Continuous Learning:** Implementing online learning mechanisms so that the model can update itself with new data over time.
- **Integration with Clinical Systems:** Seamless integration with electronic health records (EHR) to allow automated storage and retrieval of diagnostic results.
- **Regulatory Validation:** Future work should include rigorous clinical trials and validation to ensure compliance with healthcare regulations and safe deployment in real-world scenarios.

G. Summary

In summary, the future work for DermAssist AI aims to make the system more accessible, accurate, and interpretable. By combining mobile deployment, multi-modal data, self-supervised learning, and advanced explainability, the project can evolve into a comprehensive AI-assisted dermatology tool with real-world clinical utility, supporting both healthcare professionals and patients.

XVIII. Project Summary and Key Contributions

The DermAssist AI project aims to develop an accessible, accurate, and explainable system for automated skin lesion classification. The primary objectives were to leverage deep learning for multi-class skin lesion detection, integrate explainability techniques for model transparency, and provide a user-friendly interface for deployment.

A. Key Methodological Highlights

- **Deep Learning Model:** MobileNetV2 was used for efficient image classification, optimized using Memetic Algorithm to achieve high accuracy across multiple lesion categories.
- **Explainable AI:** Grad-CAM heatmaps were incorporated to visualize regions influencing model predictions, enhancing interpretability for clinicians and users.
- **Interactive Deployment:** Streamlit interface enabled easy image upload, real-time predictions, and heatmap visualization, making the system accessible to both medical professionals and general users.

B. Key Contributions

- Demonstrated that a lightweight, optimized model can achieve high accuracy while remaining computationally efficient.
- Provided an explainable framework for dermatology AI, improving trust and transparency in clinical applications.
- Developed a deployable prototype that can be extended for telemedicine, mobile applications, and educational purposes.
- Highlighted the importance of multi-modal integration, dataset diversity, and interpretability for future AI-based dermatology solutions.

C. Practical Impact

This project showcases how artificial intelligence can be translated from research to real-world application, enabling early skin lesion screening, supporting dermatologists in decision-making, and providing educational insights. The DermAssist AI framework provides a scalable and adaptable foundation for future enhancements and clinical deployment.

ACKNOWLEDGEMENT

We would like to express our sincere gratitude to all those who have contributed to the successful completion of this project.

First and foremost, we would like to thank our instructor , **[Mr Yogendra Kumar]**, for their invaluable guidance, constant encouragement, and insightful suggestions throughout the course of this work. Their expertise and support were instrumental in shaping the project and helping us overcome challenges during the development and deployment of the DermAssist AI system.

We would also like to extend our gratitude to the faculty and staff of **[SCSET/Bennett University]** for providing the necessary resources, infrastructure, and an enabling environment for research and learning.

Special thanks to our peers and colleagues for their constructive discussions, feedback, and support, which helped improve the quality and depth of our work.

Finally, we would like to acknowledge the creators of the **HAM10000 dataset** and the open-source communities for providing tools, libraries, and resources such as PyTorch, NumPy, OpenCV, Matplotlib, and Streamlit, which were essential in implementing this project.

This project would not have been possible without the combined support and encouragement of everyone mentioned above.

REFERENCES

[1] Tschandl, P., Rosendahl, C., & Kittler, H. (2018). *The HAM10000 Dataset: A Large Collection of Multi-Source Dermatoscopic Images of Common Pigmented Skin Lesions*. Scientific Data, 5, 180161. https://doi.org/10.1038/sdata.2018.161

[2] Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., & Chen, L. C. (2018). *MobileNetV2: Inverted Residuals and Linear Bottlenecks*. Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 4510–4520.

[3] Moscato, P. (1989). *On Evolution, Search, Optimization, Genetic Algorithms and Martial Arts: Towards Memetic Algorithms*. Technical Report Caltech Concurrent Computation Program, Report 826.

[4] Selvaraju, R. R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., & Batra, D. (2017). *Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization*. Proceedings of the IEEE International Conference on Computer Vision (ICCV), 618–626.

[5] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., ... & Chintala, S. (2019). *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. Advances in Neural Information Processing Systems, 32.

[6] Bradski, G. (2000). *The OpenCV Library*. Dr. Dobb's Journal of Software Tools.

[7] Harris, C. R., Millman, K. J., van der Walt, S. J., Gommers, R., Virtanen, P., Cournapeau, D., ... & Oliphant, T. E. (2020). *Array Programming with NumPy*. Nature, 585, 357–362.

[8] Hunter, J. D. (2007). *Matplotlib: A 2D Graphics Environment*. Computing in Science & Engineering, 9(3), 90–95.

[9] Streamlit Inc. (2019). *Streamlit: Turn Data Scripts into Shareable Web Apps in Minutes*. https://www.streamlit.io

[10] Esteva, A., Kuprel, B., Novoa, R. A., Ko, J., Swetter, S. M., Blau, H. M., & Thrun, S. (2017). *Dermatologist-Level Classification of Skin Cancer with Deep Neural Networks*. Nature, 542, 115–118.

[11] Pan, S. J., & Yang, Q. (2010). *A Survey on Transfer Learning*. IEEE Transactions on Knowledge and Data Engineering, 22(10), 1345–1359.

[12] Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. (2020). *A Simple Framework for Contrastive Learning of Visual Representations*. International Conference on Machine Learning (ICML), 1597–1607.

[13] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). *"Why Should I Trust You?": Explaining the Predictions of Any Classifier*. Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 1135–1144.

[14] Lundberg, S. M., & Lee, S. I. (2017). *A Unified Approach to Interpreting Model Predictions*. Advances in Neural Information Processing Systems, 30.

[15] He, K., Zhang, X., Ren, S., & Sun, J. (2016). *Deep Residual Learning for Image Recognition*. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 770–778.

[16] Huang, G., Liu, Z., Van Der Maaten, L., & Weinberger, K. Q. (2017). *Densely Connected Convolutional Networks*. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 4700–4708.

[17] Wootton, R. (2001). *Recent Advances: Telemedicine*. BMJ, 323(7312), 557–560.

[18] Rajkomar, A., Dean, J., & Kohane, I. (2019). *Machine Learning in Medicine*. New England Journal of Medicine, 380, 1347–1358.

[19] Shorten, C., & Khoshgoftaar, T. M. (2019). *A Survey on Image Data Augmentation for Deep Learning*. Journal of Big Data, 6(1), 60.

[20] Litjens, G., Kooi, T., Bejnordi, B. E., Setio, A. A. A., Ciompi, F., Ghafoorian, M., ... & van der Laak, J. A. W. M. (2017). *A Survey on Deep Learning in Medical Image Analysis*. Medical Image Analysis, 42, 60–88.